

# An information-seeking account of eye movements during spoken and signed language comprehension

Kyle MacDonald<sup>1</sup> (kylem4@stanford.edu), Aviva Blonder<sup>2</sup> (aviva.blonder@oberlin.edu),  
Virginia Marchman<sup>1</sup> (marchman@stanford.edu), Anne Fernald<sup>1</sup> (afernald@stanford.edu),  
Michael C. Frank<sup>1</sup> (mcfrank@stanford.edu)

<sup>1</sup> Department of Psychology Stanford University, <sup>2</sup> Department of Psychology Oberlin University

## Abstract

Language comprehension involves integrating information from the linguistic signal and the surrounding world. But how do we allocate our limited cognitive resources to gather different kinds of information? In the current work, we propose that the language comprehension system adapts to different tradeoffs between the value of gathering linguistic information and the value of fixating on the nonlinguistic visual world. We use two case studies of eye movements during language comprehension to test predictions of our account. First, we show that, compared to children processing spoken language, young visual language learners (a) delayed eye movements away from a language source, (b) were more accurate with their initial shifts, and (c) produced a smaller proportion of nonlanguage-driven, exploratory gaze shifts (Experiment 1). Next, we present a well-controlled, confirmatory test of our account, showing that English-speaking adults produced fewer nonlanguage-driven eye movements when processing displays of printed text compared to spoken language (Experiment 2). Together, these data suggest that the language comprehension system responds to changes in the value of seeking different kinds of information to increase the chance of rapid and accurate language understanding.

**Keywords:** eye movements; language processing; information-seeking; American Sign Language; drift diffusion models

## Introduction

Language comprehension involves the rapid integration of linguistic information with information about the surrounding world.<sup>1</sup> For example, imagine there is only one red object amongst many others, and you hear someone say, “Pass me the red \_\_\_\_!” If you have successfully encoded the visual world, then the adjective “red” allows you to constrain the speaker’s intended meaning and respond rapidly and accurately. However, for this integration process to work, we must monitor multiple streams of information in real-time using a limited set of cognitive resources. So how do we decide when to gather information about language and when to gather information about the world? We propose that the language comprehension system modulates behavior in response to changes in the value of gathering different kinds of information. We test our information-seeking account using two case studies that manipulate the value of eye movements during language processing: a) a comparison of signed and spoken language processing and b) a comparison of spoken language processing and the processing of printed text.

<sup>1</sup>Here we focus on the task of reference resolution where an interlocutor produces an utterance that refers to a concrete object in the visual scene.

The study of eye movements during language comprehension has provided insight into the interaction between conceptual representations of the world and the incoming linguistic signal. For example, work on spoken language processing shows that adults will rapidly shift visual attention upon hearing the name of an object in the visual scene, with a high proportion of these shifts occurring prior to the offset of the target word (Allopenna, Magnuson, & Tanenhaus, 1998; Tanenhaus, Spivey-Knowlton, Eberhard, & Sedivy, 1995). Moreover, researchers have found that conceptual representations activated by the visual world can modulate subsequent eye movements during language processing (Altmann & Kamide, 2007). The majority of this work has leveraged eye movements as an index of the underlying language comprehension process but has focused less on the top-down influence of participants’ goals on fixation patterns (see Salverda et al., 2011 for a review).

In contrast, researchers in the fields of natural vision and decision-making have modeled eye movements as a tool for information gathering (Ballard & Hayhoe, 2009) and a reflection of the underlying components of the decision-making process (Shadlen & Kiani, 2013). In this approach, the participant’s goal becomes critical for understanding the timing and location of fixations, and the function of eye movements is to reduce uncertainty about the external world to maximize the expected reward of future actions. For example, Hayhoe & Ballard (2005) review evidence that people do not fixate on the most salient aspects of the visual scene, but instead focus on aspects that are most helpful for the current task such as gaze shifts to an upcoming obstacle when walking.

In the current work, we leverage aspects of the information-seeking framework to account for a wider variety of fixation patterns during language comprehension. We characterize eye movements as a tradeoff between gathering more information about the visual world and gathering more information about the incoming linguistic signal. We assume that the goal of the language comprehension system is to generate behaviors that maximize the probability of making the correct future response (in the context of our task, resolving reference rapidly and accurately). To test the predictions of our account, we present two case studies where multiple kinds of information must be gathered using the same input mechanism – vision: a) processing a visual-manual language, e.g., American Sign Language (ASL), and b) processing displays of printed text. Our key prediction is that competition

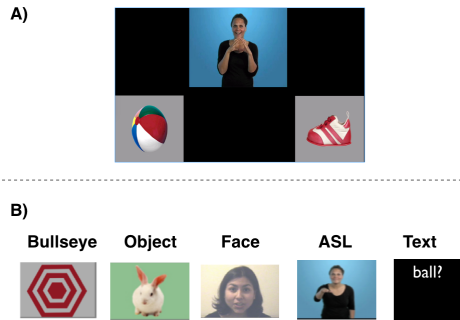


Figure 1: Stimuli information for Experiments 1 and 2. Panel A shows the layout of the three potential fixation locations across all the tasks: the center stimulus, the target picture, and the distracter picture. Panel B shows the five different center stimulus items from Experiments 1 and 2: a person signing (ASL), a person speaking (Face), a static image of a familiar object (Object), a static geometric shape (Bullseye), and printed text (Text).

for visual attention would make eye movements to the visual world less valuable, leading people to gather more information about the linguistic signal and generate a higher proportion of language-driven gaze shifts.

## Experiment 1

Experiment 1 provides an initial test of our information-seeking account. We directly compared eye movements of children learning a visual language to children learning a spoken language using two parallel, 3-alternative forced choice real-time language comprehension tasks. We predicted that the task of processing a visual language would increase the value of fixating on the language source and decrease the value of generating exploratory, nonlanguage-driven shifts to the visual world. To test this prediction, we use a combination of behavioral and model-based analyses.

## Method

**Participants** *ASL sample.* Participants were 27 native, monolingual ASL-learning children (14 deaf, 13 hearing,  $M_{Age} = 28.5$  months, range = 16-53 months). The sample consisted of both deaf children of deaf adults and hearing children of deaf adults (CODAs). All children, regardless of

Task	Mean_Age	Min_Age	Max_Age	n
ASL	27.60	16	53	27
Face	26.00	25	26	24
Object	31.90	26	39	40
Bullseye	26.10	26	27	16

Table 1: Age distributions of child participants in Experiment 1.

hearing status, were exposed to ASL from birth through extensive interaction with at least one caregiver fluent in ASL and were reported to experience at least 80% ASL in their daily lives. We excluded a small number of participants because they did not complete the task due to inattentiveness or parental interference ( $n = 5$ ).

*Spoken English samples.* Participants were 80 native, monolingual English-learning children divided across three different samples (age range = 25-39 months). Participants had no reported history of developmental or language delay. All participants were judged to be primarily English-language learners, based on parental report on linguistic input. Several participants were excluded from each of the task for not completing the task because of inattentiveness (Face:  $n = 1$ ; Bull:  $n = 5$ ; Object:  $n = \text{TODO}$ ).

Table 1 contains details about the age distributions of children in all of four tasks. We would like to note that the ASL sample included a wider age range compared to the spoken English samples. This is because native ASL learners are a difficult population to recruit, with the incidence of deafness at birth in the US being less than .003%, and only 10% of the 2-3 per 1000 children born with hearing loss have a deaf parent who is likely to be fluent in ASL (Mitchell & Karchmer, 2004). However, we included age as a covariate in all of our group comparisons and report any age effects.

**Stimuli** *ASL linguistic stimuli.* We recorded two separate sets of ASL stimuli. Both sets were recorded with a native ASL signer, using alternative ASL sentence structures for asking questions: 1) Sentence-initial wh-phrase: “HEY! WHERE [target noun]?” and 2) Sentence-final wh-phrase: “HEY! [target noun] WHERE?” To prepare the stimuli, two female native ASL users recorded several tokens of each sentence in a child-directed register. Before each sentence, the signer produced a hand-wave gesture commonly used in ASL to gain an interlocutor’s attention before initiating an utterance. The final tokens were chosen based on naturalness.

*ASL visual stimuli.* The image set consisted of colorful digitized pictures of objects presented in fixed pairs, in which the object names had no phonological overlap (cat—bird, car—book, bear—doll, ball—shoe). Images were digitized pictures presented in fixed pairs, matched for visual salience with 3–4 tokens of each object type. Each object served as target four times and as distracter four times. Side of target picture was counterbalanced across trials.

*English linguistic stimuli.* All three tasks (Object, Bullseye, and Face) featured the same female native speaker of English. The speaker used natural child-directed speech and recorded sentences with the following structure: “Look! Where’s the (target word)?”<sup>2</sup> The target words consisted of six familiar labels to match the six familiar objects: ball, banana, book, cookie, juice, and shoe.

<sup>2</sup>We would like to point out that even though there are significant differences between ASL and English question structures, all stimulus sets had the same trial structure: language to attract participants’ attention followed by a sentence containing a target noun.




Figure 2: First shift accuracy and Reaction times (RT) from Experiment 1. The left panel shows a boxplot representing the distribution of RTs for correct (orange) and incorrect (blue) shifts for each center stimulus type. The right panel shows the distribution of mean first shift accuracy scores for each center stimulus type. The solid lines represents median values, the boundaries of the box show the upper and lower quartiles, and the whiskers show the full range of the data excluding outliers.

*English video stimuli.* For the Face task, we recorded the female native English speaker as she looked straight ahead and said, “Look! Where’s the (target word)?” This video served as both a center stimulus and the source of the speech information. The video was the same size as the pictures. One video was prepared for each of the six objects. The final videos were chosen based on naturalness.

*English visual stimuli.* Features of the image set were similar to the ASL task. The image set consisted of six objects corresponding to words judged to be highly familiar to children of this age: ball, banana, book, cookie, juice, and shoes. 3–4 tokens of each of the six objects appeared throughout the experiment and side of target picture was counterbalanced across trials.

**Design and procedure** Participants viewed the ASL task on a 27” monitor. Children sat on their caregiver’s lap, and the child’s gaze was recorded using a digital camcorder set up behind the monitor. On each trial, pictures of two familiar objects appeared on the screen, a target object corresponding to the target noun, and a distracter object matched for visual salience. Between the two pictures was a central video of an adult female signing the name of one of the pictures. Participants saw 32 test trials with five filler trials (e.g. “YOU LIKE PICTURES? MORE WANT?”) interspersed to maintain children’s interest.

Participants viewed the Face, Object, and Bullseye tasks on a large projector screen in a sound-treated testing booth. Similar to the ASL task, at the beginning of each test trial, pictures of two familiar objects appeared on the screen and then a center stimulus appeared between the two pictures. The center stimulus varied across the three tasks: Face, Object, and Bullseye (see Figure 1 for details). Participants saw approximately 32 test trials with several filler trials interspersed to maintain children’s interest.

*Trial structure.* On each trial, the child saw two images of familiar objects on the screen for two seconds before the center stimulus (a signer) appeared. This time allowed the child to visually explore both images. Next, the target sentence –

which consisted of a carrier phrase, target noun, and question sign – was presented, followed by two seconds without language to allow the child to respond to the signer’s sentence. The trial structure of the Face, Object, and Bullseye tasks were highly similar: children were given two seconds to visually explore the objects prior to the appearance of the center stimulus, then processed a target sentence, and finally were given two seconds of silence to generate a response to the target noun.

*Coding.* Participants’ gaze patterns were videotaped and later coded frame-by-frame at 33-ms resolution by trained coders blind to target side. On each trial, coders indicated whether the eyes were fixated on the central signer, one of the images, shifting between pictures, or away (off), yielding a high-resolution record of eye movements aligned with target noun onset. Prior to coding, all trials were pre-screened to exclude those few trials on which the participant was inattentive or there was external interference.

**Behavioral measures** *Reaction time.* Reaction time (RT) corresponds to the latency to shift from the central stimulus to either the target or the distracter pictures measured from target-noun onset. We chose to cutoff RTs longer than two seconds since these shifts are unlikely to be generated in response to the incoming language stimulus (see Ratcliff, 1993).

*First shift accuracy.* Accuracy was the mean proportion of first shifts that landed on the target picture out of the total number of shifts that landed on either the target or the distracter picture over a two-second window from target noun onset. This measure reflects the accuracy of language-driven saccades to the visual world. Mean first shift accuracy scores were computed for each participant for both correct and incorrect shifts. Trials where the participants did not generate a shift were not included in the computation.

## Results and Discussion

**Analysis plan** First, we present behavioral analyses of accuracy and RT. Since RTs are not normally distributed, we

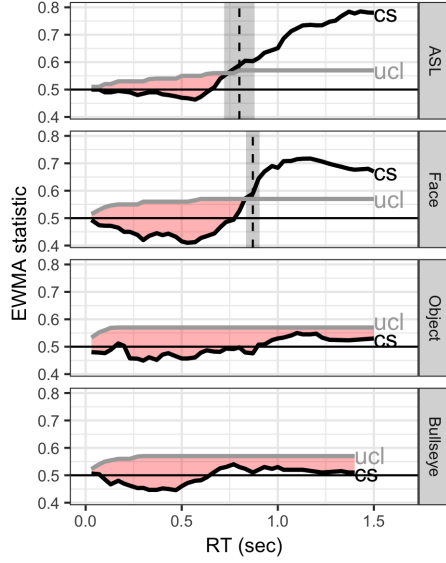


Figure 3: Output for the EWMA guessing model for all center stimulus types in Experiment 1. The black curve represents the evolution of the control statistic (cs) as a function of reaction time. The grey curve represents the upper control limit (ucl). The vertical dashed line is the median cutoff value (point when the control process shifts out of a guessing state) across all participants. The grey shaded area represents the 95% confidence interval around the estimate of the median cutoff point. And the shaded red area represents the proportion of responses that were flagged as guesses by the EWMA model.

log transformed all RTs for our statistical analyses. To quantify differences between groups, we used the `lme4` R package (Bates, Maechler, Bolker, & Walker, 2013) to fit mixed-effects regression models that included a by-subject random intercept to account for repeated measures from each participant. All data and analysis code can be found in the online repository for this project: <https://github.com/kemacdonald/speed-acc>.

Next, we present two model-based analyses that quantify different patterns of eye movements across the four language comprehension tasks. First, we use an Exponentially weighted moving average (EWMA) method (Vandekerckhove & Tuerlinckx, 2007) to model the proportion of nonlanguage-driven shifts away from the center stimulus to the visual world. The goal of the EWMA is to identify whether a process has deviated from a pre-defined “control state” by taking into account the prior behavior of the process, weighting recent observations more heavily. The model generates two values: a “control statistic” (CS) and an “upper control limit” (UCL) for each point in the RT distribution. Once the CS becomes larger than the UCL, the process is determined to have exited the control state.<sup>3</sup>

<sup>3</sup> $c_s = \lambda x_s + (1 - \lambda)c_{s-1}$  where the  $\lambda$  parameter determines the number of prior RTs that are included in the moving average compu-




Figure 4: Posterior distributions over the boundary and drift rate parameters in the hierarchical drift diffusion model.

Here, we adapt the EWMA approach to model changes in the process that generate eye movements away from the center stimulus to the visual world. We define the control state as an expectation of nonlanguage-driven shifts and model this as a Bernoulli process with probability of success 0.5. As the sentence unfolds, we assume that participants gather more of the linguistic information prior to shifting and the underlying process should bias towards more accurate shifts or a Bernoulli process with probability success  $> 0.5$ . With this model, we can compare across our groups: a) the cutoff point when the CS exceeded the UCL indicating that participants started to generate language-driven shifts and b) the proportion of shifts that the model categorizes as language-driven vs. exploratory.

Finally, we use Drift diffusion models (DDMs) (Ratcliff & Childers, 2015) to quantify differences in the dynamics of speed and accuracy for eye movements generated in response to the incoming linguistic signal. DDMs form a class of sequential decision-making models designed specifically for rapid two-alternative forced choice tasks. The model assumes that people accumulate noisy evidence in favor of one alternative with a response generated when the evidence crosses a pre-defined decision threshold. The DDM approach is useful because it can account for all components of the behavioral response: correct and incorrect RT distributions. Moreover, the parameters of the DDM map onto meaningful psychological variables of interest. Here we focus on two of these parameters: **boundary separation**, which maps onto the amount of evidence gathered before generating a response (higher values suggest a prioritization of accuracy over speed) and **drift rate**, which maps onto the amount of evidence that is accumulated per unit time (higher values indicate more efficient processing).

We chose to implement a hierarchical Bayesian version of the DDM using the HDDM Python package (Wiecki, Sofer,

tation.  $UCL_s = c_0 + L\sigma_0\sqrt{\frac{\lambda}{2-\lambda}[1 - (1-\lambda)^{2s}]}$  where  $L$  controls the width of the control limits with higher values leading to a more conservative test. We chose values used in prior work using the EWMA with speeded decision tasks (Vandekerckhove & Tuerlinckx, 2007)

& Frank, 2013) because we were dealing with relatively few trials from child participants and recent simulation studies have shown that the HDDM approach was better than other DDM fitting methods when the number of observations was small (Ratcliff & Childers, 2015).

**Behavioral analyses** *RT.* Panel A of Figure 2 shows children’s RTs for correct and incorrect shifts. Visual inspection of the figure suggests that there was a speed accuracy trade-off in the ASL, Face, and Bullseye conditions where incorrect RTs tended to be faster than correct RTs. To quantify differences across the groups, we fit a linear mixed-effects regression predicting first shift RT as a function of center stimulus type, controlling for age, and including user-defined contrasts to test specific comparisons of interest:  $\text{Log}(\text{RT}) \sim \text{center stimulus type} + \text{age} + (1 \mid \text{subject})$ . We found that (a) ASL learners delayed their shifts compared to all of the spoken English participants ( $\beta = -0.96$ ,  $p < .001$ ), (b) ASL learners’ shifts were slower compared to participants in the Face task ( $\beta = -0.42$ ,  $p < .001$ ), and (c) participants in the Face task shifted slower compared to participants in the Object and Bullseye tasks ( $\beta = -0.73$ ,  $p < .001$ ).

*Accuracy.* Next we compared the accuracy of first shifts across the different tasks by fitting a mixed-effects logistic regression with the same specifications and contrasts as the RT model. We found that (a) ASL learners were more accurate compared to all of the spoken English participants ( $\beta = -0.69$ ,  $p < .001$ ), (b) ASL learners were more accurate when directly compared to participants in the Face task ( $\beta = -0.54$ ,  $p = 0.001$ ), and (c) participants in the Face task were numerically more accurate compared to participants in the Object and Bullseye tasks ( $\beta = -0.73$ ) but this effect was not significant at the .05 level ( $p = 0.094$ ).

**Model-based analyses** *EWMA.* Figure 3 shows the output of the EWMA model for the four different center stimulus types. This plot shows changes in the control statistic (CS) and the upper control limit (UCL) as a function of participants’ reaction times. Each CS starts at chance performance and below the UCL. In the ASL and Face tasks, the CS values begin to increase with RTs around 0.7 seconds after noun onset eventually exceeding the UCL and entering a nonguessing state. In contrast, the CS in the Object and Bullseye tasks never crosses the UCL, meaning that children’s first shifts were equally likely to land on the target as on the distracter, regardless of when those shifts were initiated in the RT distribution. This result suggests that first shifts in the Bullseye/Object tasks are language-driven and may instead reflect a different process such as gathering more information about the referents in the visual world.

Next, we compared the EWMA output of participants in the ASL and Face tasks. We found that ASL learners generated fewer shifts when the CS was below the UCL ( $\beta = -1.64$ ,  $p < .001$ ), indicating that a larger proportion of their initial shifts away from the language source were driven by the incoming linguistic information (see the differences in the red

shaded area in Figure 3). We did not find evidence for a difference in the location of the point in the RT distribution at which the CS crossed the UCL ( $\beta = -0.03$ ,  $p = 0.576$ ). The EWMA analyses provide evidence that, compared to children processing spoken language, ASL learners shifts were more likely to be language-driven.

*DDM.* Using the output of the EWMA, we could now compare the timing and accuracy of responses generated by the incoming linguistic signal by participants the ASL and Face tasks. Figure 4 shows the full posterior distributions over the boundary separation and drift rate parameters for the ASL and Face tasks. We found that ASL learners had a higher estimate for the boundary separation parameter compared to the Face participants (ASL boundary = 1.77, HDI = [1.64, 1.9]; Face boundary = 1.35, HDI = [1.21, 1.49]), with no overlap in the credible values. This suggests that ASL learners accumulated more evidence about the linguistic signal before generating an eye movement. We did not find any difference in the drift rate parameter, indicating that both groups processed the linguistic information with similar efficiency (ASL drift = 0.64, HDI = [0.44, 0.83]; Face drift = 0.57, HDI = [0.33, 0.83]).

Taken together, the behavioral analyses and the EWMA/HDDM results provide converging support that ASL learners were sensitive to the different value of eye movements, producing fewer nonlanguage-driven shifts to explore the visual world and prioritizing accuracy over speed when generating a saccade away from a language source. This behavior seems reasonable since the potential for missing subsequent linguistic information is high if ASL users shifted prior to gathering sufficient information. It is important to point out that these findings were fundamentally exploratory and that there were several, potentially important differences between the stimuli, apparatus, and populations. Thus, we set out to perform a well-controlled, confirmatory test in Experiment 2.

## Experiment 2

The goal of Experiment 2 was to replicate a key finding from Experiment 1: that increasing the competition between attention language and attention to the visual world reduces the proportion of nonlanguage-driven shifts. Moreover, we set out to conduct a confirmatory test of our account that also controlled for the population differences in Experiment 1. To accomplish these goals, we tested a sample of English-speaking adults using a within-participants manipulation of the center stimulus type. We used the Face and Bullseye stimulus sets from Experiment 1<sup>4</sup> and added two new conditions: Text, where the verbal language information was accompanied by an unfolding display of printed text, and Text-no-audio, where the spoken language stimuli were removed (see Figure 1 for an example of the text display). The key behavioral prediction is that participants in the Text condi-

<sup>4</sup>We chose not to include the Object stimulus set because performance between the Bullseye and Object tasks was indistinguishable in Experiment 1.



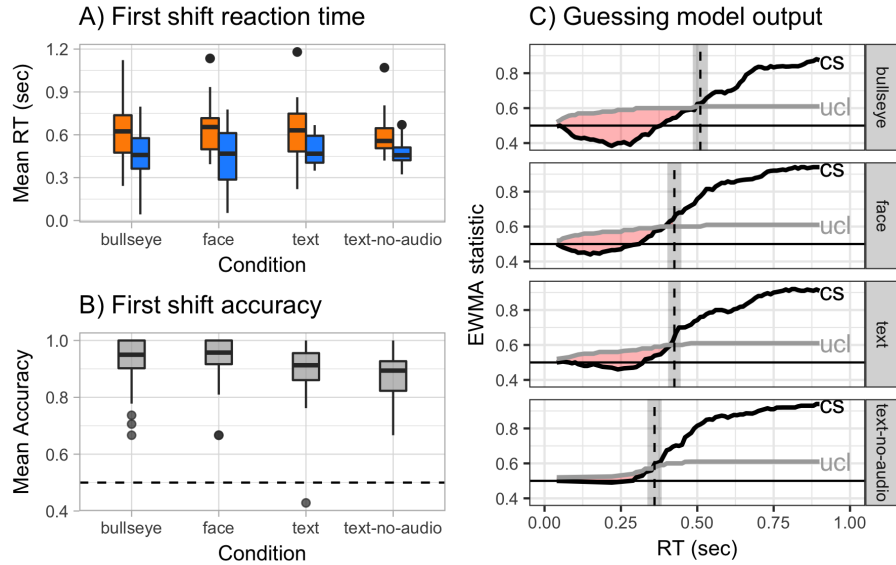


Figure 5: Behavioral and model-based results from Experiment 2 for all center stimulus types. Panel A shows reaction times, Panel B shows accuracy, and Panel C shows the output of the EWMA guessing model. All plotting conventions are the same as in Figures 2 and 3.

tions (where the center stimulus provides information about the linguistic stimulus) should produce a higher proportion of language-driven shifts and fewer exploratory shifts to the visual world. We did not predict a difference in the DDM analysis since adults were responding to highly familiar words with only two response options.

## Method

**Participants** 25 Stanford undergraduates participated (6 male, 20 females) for course credit. All participants were monolingual, native English speakers and had normal vision.

**Stimuli** Audio and visual stimuli were identical to the Face and Bullseye tasks in Experiment 1. We included a new center fixation stimulus type: printed text. The text was displayed in a white font on a black background. The text was programmed such that only a single word appeared on the screen at a time with each word appearing for the same duration as the corresponding word in the spoken language stimuli.

**Design and procedure** The design was nearly identical to Experiment 1, with the exception of the change to a within-subjects manipulation where each participant completed all four tasks (Bullseye, Face, Text, and Text-no-audio). In both the Text and Text-no-audio conditions, the center stimulus was white text printed on a black background. In the Text condition, spoken language accompanied the printed text. In the Text-no-audio condition, the spoken language stimuli was removed. Participants viewed each task on a 27" monitor and were told their eye movements would be recorded. Participants saw four blocks of 32 test trials for a total of 128 trials. Participants' eye movements were tracked using automated eye-tracking software from SensoMotoric Instruments (SMI),

sampling at 120 Hz.

## Results and Discussion

**Behavioral analyses RT.** Panel A of Figure 5 shows adults' RTs for correct and incorrect shifts. Visual inspection of the figure suggests that there was a speed-accuracy tradeoff in all conditions: incorrect RTs tended to be faster than correct RTs. We fit a linear mixed-effects regression with the same specification as in Experiment 1, but we added by-subject intercepts and slopes for each center stimulus type to account for our within-subjects condition manipulation. We did not find evidence that RTs were different across conditions (all  $p > .05$ ).

**Accuracy.** Next, we compared the accuracy of first shifts across the different tasks by fitting a mixed-effects logistic regression with the same specifications (see Panel B of Figure 5). We found that participants tended to be less accurate in the Text conditions compared to conditions without text ( $\beta = -0.62$ ,  $p < .001$ ). We did not find evidence for any of the other comparisons.

**Model-based analyses EWMA** Panel C of Figure 5 shows the output of the EWMA model for the four different center stimulus types. In contrast to Experiment 1, we found that the CS crosses the UCL in all four conditions, suggesting that at some point in the RT distribution adults' shifts are driven by seeking out the named referent. Interestingly, we found a graded effect of condition on the location when the CS crossed the UCL such that the Text-no-audio condition occurred first (see the vertical dashed lines in Figure 5), followed by the Text and Face conditions that were not different from one another, and finally the Bullseye condition

(TODO stats). We also found the same graded difference in the proportion of shifts that occurred while the CS was below the UCL, indicating a higher proportion of first shifts were language-driven in the Text conditions (TODO stats). These results provide strong evidence for our key prediction: that increasing the value of fixating the center stimulus for gathering language information reduces the probability of eye movements to gather information about the visual world.

**DDM.** Using the output of the EWMA, we fit the same hierarchical DDM as in Experiment 1 to test if there were differences in the timing and accuracy of responses generated by the incoming linguistic signal. There was a high overlap of the posterior distributions for both the boundary separation and the drift rate parameters, providing evidence that the estimates did not differ across conditions.

Short summary of E2 results here.

## General Discussion

Language comprehension involves a tradeoff between gathering linguistic information and gathering information about the surrounding world. In the current work, we propose that eye movements during language processing reflect a sensitivity to the value of gathering different kinds of information. We found that, young ASL-learners generated slower but more accurate shifts away from a language source and produced a smaller proportion of nonlanguage-driven shifts. We found the same pattern of behavior within a sample of English-speaking adults processing displays of printed text compared to spoken language. These results suggest that as the value of fixating on a location to gathering information about language increases, eye movements to gather information about the visual world become less useful and occur less frequently.

There are several limitations to our experiments. First, we have not performed a confirmatory test of the DDM findings (ASL-learners prioritization of accuracy over speed) from Experiment 1. So this finding, while interesting, is preliminary. Second, we do not know what might be driving the population differences in Experiment 1. It could be that ASL-learners massive experience dealing with competition for visual attention leads to changes in the deployment of eye movements during language comprehension. Or, it could be that the in-the-moment constraints of processing a visual language causes different fixation behavior. Finally, we used a very simple visual world, with only three places to look, and very simple linguistic stimuli, especially for the adults in Experiment 2. Thus it remains an open question how these results might scale up to more complex language information and visual environments.

A strength of this work is the attempt to integrate top-down, goal-based models of vision with work on language-driven eye movements. While we chose to start with two particular case studies – ASL and text processing – we think the account is more general and that there are many real world situations where people must negotiate the tradeoff between gathering

more information about language or about the world: e.g., processing spoken language in noisy environments or at a distance; or early in language learning when children are acquiring new words and often rely on nonlinguistic cues to reference such as pointing or eye gaze. Overall, we hope this work contributes to a broader account of eye movements during language comprehension that can explain fixation behaviors across wider variety of populations and processing contexts.

## Acknowledgements

We are grateful to the families who participated in this research and to the California School for the Deaf in Fremont for their help with participant recruitment. This work was supported by a National Science Foundation Graduate Research Fellowship to KM and a NIDCD grant to Anne Fernald and David Corina (DC012505).

## References

- Allopenna, P. D., Magnuson, J. S., & Tanenhaus, M. K. (1998). Tracking the time course of spoken word recognition using eye movements: Evidence for continuous mapping models. *Journal of Memory and Language*, 38(4), 419–439.
- Altmann, G., & Kamide, Y. (2007). The real-time mediation of visual attention by language and world knowledge: Linking anticipatory (and other) eye movements to linguistic processing. *Journal of Memory and Language*, 57(4), 502–518.
- Ballard, D., & Hayhoe, M. (2009). Modelling the role of task in the control of gaze. *Visual Cognition*, 17(6-7), 1185–1204.
- Bates, D., Maechler, M., Bolker, B., & Walker, S. (2013). lme4: Linear mixed-effects models using eigen and s4. r package version 1.0-5.
- Hayhoe, M., & Ballard, D. (2005). Eye movements in natural behavior. *Trends in Cognitive Sciences*, 9(4), 188–194.
- Ratcliff, R., & Childers, R. (2015). Individual differences and fitting methods for the two-choice diffusion model of decision making. *Decision*, 2(4), 237–279.
- Salverda, A. P., Brown, M., & Tanenhaus, M. K. (2011). A goal-based perspective on eye movements in visual world studies. *Acta Psychologica*, 137(2), 172–180.
- Shadlen, M. N., & Kiani, R. (2013). Decision making as a window on cognition. *Neuron*, 80(3), 791–806.
- Tanenhaus, M. K., Spivey-Knowlton, M. J., Eberhard, K. M., & Sedivy, J. C. (1995). Integration of visual and linguistic information in spoken language comprehension. *Science*, 268(5217), 1632.
- Vandekerckhove, J., & Tuerlinckx, F. (2007). Fitting the ratcliff diffusion model to experimental data. *Psychonomic Bulletin & Review*, 14(6), 1011–1026.
- Wiecki, T. V., Sofer, I., & Frank, M. J. (2013). HDDM: Hierarchical bayesian estimation of the drift-diffusion model in python. *Frontiers in Neuroinformatics*, 7, 14.